

# CS202 Lab 2 Report

Yingfan Gao  
ygao195@ucr.edu  
UC Riverside

Yufei Li  
yli927@ucr.edu  
UC Riverside

Shahab Nikkhoo  
snikk002@ucr.edu  
UC Riverside

## 1 THE LIST OF ALL FILES MODIFIED

In this assignment, we have modified the following files:

- kernel/syscall.h: Define new system call numbers.
- kernel/syscall.c: Update system call table.
- kernel/sysproc.c: Define system call function.
- kernel/proc.c: Create new kernel function.
- kernel/proc.h: Create new process states such as stride, tickets for scheduling.
- kernel/defs.h: Create new system call.
- user/usys.pl and user/user.h: Update userspace system call interface.
- user/lab2.c: Write a test program.
- Makefile: Lab2 operating with different scheduling algorithms.

## 2 EXPLANATION ON WHAT CHANGES WE MADE

Similar to the previous assignment, we first define new system call #25 [SYS\_sched\_statistics] and #26 [SYS\_sched\_tickets] (see Figure 1), and then update the corresponding system call table (see Figure 2).

```
21 #define SYS_mkdir 20
22 #define SYS_close 21
23 #define SYS_hello 22
24 #define SYS_sysinfo 23 //lab1
25 #define SYS_procinfo 24 //lab1
26 #define SYS_sched_statistics 25 //lab2
27 #define SYS_sched_tickets 26 //lab2
```

Figure 1: Syscall.h

```
134 [SYS_close] sys_close,
135 [SYS_hello] sys_hello,
136 [SYS_sysinfo] sys_info, // lab1
137 [SYS_procinfo] sys_procinfo, // lab1
138 [SYS_sched_statistics] sys_sched_statistics, // lab2
139 [SYS_sched_tickets] sys_sched_tickets, // lab2
```

Figure 2: Syscall.c

Next, we defined system call function [sys\_sched\_statistics(void)] as well as [sys\_sched\_tickets(void)] and program it in "kernel/proc.c" (as shown in Figure 3, Figure 4 and Figure 5).

```
122 uint64
123 sys_sched_statistics(void) //lab2
124 {
125     sched_statistics();
126     return 0;
127 }
128
129
130 uint64
131 sys_sched_tickets(void) //lab2
132 {
133     int tickets;
134     // int ticks;
135     argint(0, &tickets);
136     // argint(0, &ticks);
137     sched_tickets(tickets);
138     return 0;
139 }
```

Figure 3: Sysproc.c

```
826 int sched_statistics(void)
827 {
828     struct proc * p = myproc();
829     // acquire(&p->lock);
830     for(p = proc; p < &proc[NPROC]; p++) {
831         if(p->state != UNUSED)
832         {
833             if(p->tickets != 60) // child process
834             {
835                 printf("%d(%s): tickets: %d, ticks: %d\n",
836                     p->pid, p->name, p->tickets, p->ticks);
837             }
838             if (p->tickets == 60) // parent process
839             {
840                 printf("%d(%s): tickets: xxx, ticks: %d\n", p->pid, p->name, p->ticks);
841             }
842         }
843     }
844     // release(&p->lock);
845     return 0;
846 }
847
848 // This system call sets the caller process's ticket value to the given parameter.
849 int sched_tickets(int tickets)
850 {
851     struct proc * p = myproc();
852     // acquire(&p->lock);
853     p->tickets = tickets;
854     #ifdef STRIDE
855     p->stride = 10000/tickets;
856     #endif
857     // release(&p->lock);
858     return 0;
859 }
```

Figure 4: Proc.c (system call)

```
89 enum procstate state; // Process state
90 void *chan; // If non-zero, sleeping on chan
91 int killed; // If non-zero, have been killed
92 int xstate; // Exit status to be returned to parent's wait
93 int pid; // Process ID
94 int tickets; // Number of tickets for lottery scheduler
95 int ticks; // Number of time slices for lottery scheduler
96 #ifdef STRIDE
97 int stride; // Stride for stride scheduler
98 int pass; // Pass for stride scheduler
99 #endif
```

Figure 5: Proc.h

Also, we add these functions into "kernel/defs.h" that enable the new created system call available. Then, we update userspace

system call interface with the new system calls (as shown in Figure 6, Figure 7, and Figure 8).

```
112 int sysinfo(int); //lab1
113 extern uint systemcallcount;
114 int procinfo(struct pinfo*); //lab1
115 int sched_statistics(void); //lab2
116 int sched_tickets(int); //lab2
117 unsigned short rand(void); //lab2
```

Figure 6: Defs.h

```
40 entry("sysinfo"); # lab1
41 entry("procinfo"); # lab1
42 entry("sched_statistics"); # lab2
43 entry("sched_tickets"); # lab2
```

Figure 7: Usys.pl

```
26 int hello(void);
27 int sysinfo(int); //lab1
28 int procinfo(struct pinfo*); //lab1
29 int sched_statistics(void); //lab2
30 int sched_tickets(int); //lab2
```

Figure 8: User.h

The two scheduling algorithms, i.e., lottery scheduling and stride scheduling are shown in Figure 9 and Figure 10, respectively.

```
471 // Do lottery scheduling
472 #ifdef LOTTERY
473 // Get total number of tickets
474 int total_tickets = 0;
475 for(p = proc; p < &proc[NPROC]; p++){
476     acquire(&p->lock);
477     if(p->state == RUNNABLE){
478         total_tickets += p->tickets;
479     }
480     release(&p->lock);
481 }
482 int winner = 1 + ((int)rand() % total_tickets);
483 int current = 0;
484 // Find the scheduled process
485 for(p = proc; p < &proc[NPROC]; p++){
486     acquire(&p->lock);
487     if(p->state != RUNNABLE) {
488         release(&p->lock);
489         continue;
490     }
491     current += p->tickets;
492     if(current < winner){
493         release(&p->lock);
494         continue;
495     }
496     p->state = RUNNING;
497     c->proc = p;
498     p->ticks += 1; // Increment ticks
499     switch(&c->context, &p->context);
500     c->proc = 0;
501     release(&p->lock);
502     break;
503 }
504 #endif
```

Figure 9: Proc.c (Lottery scheduling)

For lottery scheduling, we first statistic the total number of tickets by iterating all runnable processes. Then we generate a winning score using the rand() function (provided in Part2 pseudo code). The accumulative tickets for a process, if larger than the winning score, would indicate that process is scheduled to be executed. We then add one more time slice to its ticks and implement context switch, with the lock then being released.

```
506 #ifdef STRIDE
507 // Do stride scheduling
508 int current = 100000;
509 struct proc *chosen = proc; // Pointer to the chosen process
510 // Find the scheduled process
511 for(p = proc; p < &proc[NPROC]; p++){
512     acquire(&p->lock);
513     if(p->state != RUNNABLE) {
514         release(&p->lock);
515         continue;
516     }
517     if(p->pass < current) {
518         current = p->pass;
519         chosen = p;
520     }
521     release(&p->lock);
522 }
523 acquire(&chosen->lock);
524 chosen->state = RUNNING;
525 c->proc = chosen;
526 chosen->ticks += 1; // Increment ticks
527 chosen->pass += chosen->stride; // Increment pass
528 switch(&c->context, &chosen->context);
529 c->proc = 0;
530 release(&chosen->lock);
531 #endif
```

Figure 10: Proc.c (Stride scheduling)

For stride scheduling, we use a process pointer "chosen" to keep track of all runnable processes and point at the one with minimum pass (cumulative stride). Then, we add one more time slice to that pointer and implement context switch, with the lock then being released.

### 3 EXPERIMENT FIGURES AND DISCUSSION

In this section, we represent the comparison between Lottery scheduling and Stride Scheduling. We plot lottery scheduling (Figure 11 and Figure 12), and stride scheduling (Figure 13 and Figure 14), both under two time allocations (8:4:2:1 and 1:1:1:1).

p1, p2, p3 and p4




Figure 11: Lottery scheduling (8-4-2-1)




Figure 12: Lottery scheduling (1-1-1-1)

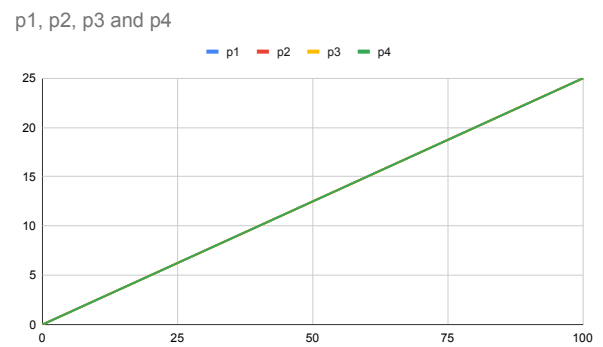


Figure 14: Stride scheduling (1-1-1-1)




Figure 13: Stride scheduling (8-4-2-1)

From the above figures, we can see that stride scheduling is much more deterministic than lottery scheduling. For each scheduling algorithm, we run the two experiments for 5 times to get the average.

#### 4 SUMMARY OF CONTRIBUTIONS

In this lab assignment, Yufei took charge of the first part of the lab, while Shahab completed the second part. And Yingfan was responsible for the writing of lab report.